

# MSAS Tutorial Sequence -Module Five-

---

WRITTEN BY ADVAY MUCHOOR

A solid yellow horizontal bar spanning the width of the slide at the bottom.

# GAME TIME



# Goals of Today

---

1. Data cleaning
2. More row operations
3. Getting information from your data
  1. Column information operations
  2. Data Visualizations
4. Ideas for the last tutorial?



# Cleaning Data

---

WHY DO WE NEED TO DO IT?

# Values Out of Range

```
# Create a copy of passing (need to use the copy function in order to create a new dataframe)
best_ratio = passing_stats.copy(deep=True)
# If we do not use the copy function, changing our new table will also change our original (which we don't want!!!)

best_ratio["TD Ratio"] = round(best_ratio["TD"] / best_ratio["Int"], 2)
best_ratio = best_ratio.sort_values("TD Ratio", ascending = False)
best_ratio.head()
```

	Rk	Player	Tm	Age	Pos	G	GS	Cmp	Att	Cmp%	Yds	TD	Int	Y/A	Y/C	Y/G	Rate	QBR	Sk	TD Ratio
105	106	Albert Wilson	MIA	26	wr	7	3	1	1	100.0	52	1	0	52.0	52.0	7.4	158.3	99.7	0	inf
74	75	Mohamed Sanu	ATL	29	WR	16	16	1	2	50.0	5	1	0	2.5	5.0	0.3	95.8	32.3	0	inf
84	85	Kevin Byard	TEN	25	FS	16	16	1	1	100.0	66	1	0	66.0	66.0	4.1	158.3	NaN	0	inf
81	82	Chris Boswell	PIT	27	K	15	0	1	1	100.0	2	1	0	2.0	2.0	0.1	118.7	NaN	0	inf
77	78	Danny Amendola	MIA	33	WR	15	15	1	1	100.0	28	1	0	28.0	28.0	1.9	158.3	100.0	0	inf

# Invalid/Unwanted Data

---

```
best_completions = passing_stats.sort_values("Cmp%", ascending = False)
best_completions.head()
```

	Rk	Player	Tm	Age	Pos	G	GS	Cmp	Att	Cmp%	Yds	TD	Int	Y/A	Y/C	Y/G	Rate	QBR	Sk
105	106	Albert Wilson	MIA	26	wr	7	3	1	1	100.0	52	1	0	52.0	52.0	7.4	158.3	99.7	0
94	95	Sam Koch	BAL	36	P	16	0	1	1	100.0	21	0	0	21.0	21.0	1.3	118.7	11.8	0
70	71	Odell Beckham	NYG	26	WR	12	12	2	2	100.0	106	2	0	53.0	53.0	8.8	158.3	NaN	0
71	72	Julian Edelman	NWE	32	WR	12	12	2	2	100.0	43	0	0	21.5	21.5	3.6	118.7	94.3	0
76	77	Nelson Agholor	PHI	25	WR	16	16	1	1	100.0	15	0	0	15.0	15.0	0.9	118.7	100.0	0

# Filtering

---

- Most of these cleaning operations can be taken care of simply by using inequalities
  - We learned this last module
    - Check if the position is a quarterback
    - Set a threshold for the number of passes thrown
    - Set a threshold for the number of games thrown/started
- We can exclude these rows from our dataset simply because we do not need them
- But what about the rows where values are undefined?

# Missing values

```
In [41]: ▶ passing_stats = pd.read_excel("nfl_stats.xlsx")  
passing_stats
```

85	86	Larik Cohen*+	CHI	23	rb/wr	16	7	1	1	100.0	1	1	0	1.0	1.0	0.1	118.7	100.0	0
86	87	Logan Cooke	JAX	23	P	16	0	1	1	100.0	4	0	0	4.0	4.0	0.3	83.3	NaN	0
87	88	Eric Ebron*	IND	25	te/wr	16	8	0	1	0.0	0	0	0	0.0	NaN	0.0	39.6	0.0	0
88	89	Bruce Ellington	2TM	27	NaN	7	3	0	1	0.0	0	0	0	0.0	NaN	0.0	39.6	NaN	0
89	90	Larry Fitzgerald	ARI	35	WR	16	16	1	1	100.0	32	1	0	32.0	32.0	2.0	158.3	100.0	0
90	91	Dontrell Hilliard	CLE	23	NaN	11	0	0	1	0.0	0	0	1	0.0	NaN	0.0	0.0	0.0	0
91	92	DeAndre Hopkins*+	HOU	26	WR	16	16	0	1	0.0	0	0	0	0.0	NaN	0.0	39.6	2.9	0
92	93	Darius Jennings	TEN	26	NaN	16	0	1	1	100.0	21	0	0	21.0	21.0	1.3	118.7	99.9	0
93	94	Zay Jones	BUF	23	WR	16	15	0	1	0.0	0	0	0	0.0	NaN	0.0	39.6	0.0	0
94	95	Sam Koch	BAL	36	P	16	0	1	1	100.0	21	0	0	21.0	21.0	1.3	118.7	11.8	0
95	96	Christian McCaffrey	CAR	22	RB	16	16	1	1	100.0	50	1	0	50.0	50.0	3.1	158.3	100.0	0
96	97	Anthony Miller	CHI	24	wr	15	4	1	1	100.0	8	0	0	8.0	8.0	0.5	100.0	81.2	0
97	98	Matt Prater	DET	34	K	16	0	1	1	100.0	8	1	0	8.0	8.0	0.5	139.6	NaN	0
98	99	Emmanuel Sanders	DEN	31	WR	12	12	1	1	100.0	28	1	0	28.0	28.0	2.3	158.3	100.0	0



# “NaN” Values

---

- An “NaN” value is a value that is undefined
- Stands for “Not a Number”
  - Usually placeholders for cells that were left empty in the original DataFrame
  - Can also be used to represent division by zero in some cases
- Why do we care about NaN values?
  - Disruptive to our DataFrame
  - Messes up operations on a series we want to do

# Why are NaN values so bad???

```
# fails because NaN values exist
all_qbs = passing_stats[passing_stats["Pos"].str.contains("qb")]

-----

--
ValueError                                Traceback (most recent call last)
<ipython-input-94-47b7ff64c0cb> in <module>
      1 # fails because NaN values exist
----> 2 all_qbs = passing_stats[passing_stats["Pos"].str.contains("qb")]

~\Anaconda3\lib\site-packages\pandas\core\frame.py in __getitem__(self, key)
    2915
    2916         # Do we have a (boolean) 1d indexer?
-> 2917         if com.is_bool_indexer(key):
    2918             return self._getitem_bool_array(key)
    2919

~\Anaconda3\lib\site-packages\pandas\core\common.py in is_bool_indexer(key)
    122         if not lib.is_bool_array(key):
    123             if isna(key).any():
--> 124                 raise ValueError(na_msg)
    125             return False
    126             return True

ValueError: cannot index with vector containing NA / NaN values
```

# Locate rows with NaN values

```
# Find the rows that have missing values
# Allows you to get a better feel for the types of errors that are occurring
nan_rows = passing_stats[passing_stats.isnull().any(axis=1)]
nan_rows.head(10)
```

	Rk	Player	Tm	Age	Pos	G	GS	Cmp	Att	Cmp%	Yds	TD	Int	Y/A	Y/C	Y/G	Rate	QBR	Sk
50	51	DeShone Kizer	GNB	22	NaN	3	0	20	42	47.6	187	0	2	4.5	9.4	62.3	40.5	25.2	4
55	56	Mike Glennon	ARI	29	NaN	2	0	15	21	71.4	174	1	0	8.3	11.6	87.0	112.0	77.7	1
56	57	Matt Cassel	DET	36	NaN	2	0	7	17	41.2	59	0	1	3.5	8.4	29.5	26.3	8.8	1
57	58	Joshua Dobbs	PIT	23	NaN	5	0	6	12	50.0	43	0	1	3.6	7.2	8.6	24.0	58.0	0
59	60	Matt Schaub	ATL	37	NaN	3	0	5	7	71.4	20	0	0	2.9	4.0	6.7	74.1	60.6	0
60	61	Robert Griffin	BAL	28	NaN	3	0	2	6	33.3	21	0	0	3.5	10.5	7.0	44.4	2.1	0
61	62	Kyle Lauletta	NYG	23	NaN	2	0	0	5	0.0	0	0	1	0.0	NaN	0.0	0.0	0.1	0
62	63	Jacoby Brissett	IND	26	NaN	4	0	2	4	50.0	2	0	0	0.5	1.0	0.5	56.2	100.0	0
63	64	Johnny Hekker	LAR	28	P	16	0	2	4	50.0	19	0	0	4.8	9.5	1.2	63.5	NaN	0
64	65	Geno Smith	LAC	28	NaN	5	0	1	4	25.0	8	0	0	2.0	8.0	1.6	39.6	0.7	1

# Fill all NaN values

---

```
# Fill all NaN values with the same value
# Look at the 50th row (as seen in the previous cell) to see the effective changes
fill_zeros = passing_stats.fillna(0)
fill_zeros.loc[[50]]
```

	Rk	Player	Tm	Age	Pos	G	GS	Cmp	Att	Cmp%	Yds	TD	Int	Y/A	Y/C	Y/G	Rate	QBR	Sk
50	51	DeShone Kizer	GNB	22	0	3	0	20	42	47.6	187	0	2	4.5	9.4	62.3	40.5	25.2	4

- When would filling all values with the same value be useful
  - If the entire table is comprised of strings, we can fill an NaN cell with an empty string
  - If the entire table is filled with integers, we can fill with a -1 and only check values that are positive
- Other than this, there aren't many use cases
  - It is better to be more specific in the ways you want to deal with NaN values
  - How can we do this?

# Option #1 – Delete all rows with NaN values

---

```
# remove all rows with NaN values
removed_NaN = passing_stats[~passing_stats.isnull().any(axis=1)]
removed_NaN.head()
```

	Rk	Player	Tm	Age	Pos	G	GS	Cmp	Att	Cmp%	Yds	TD	Int	Y/A	Y/C	Y/G	Rate	QBR	Sk
0	1	Ben Roethlisberger	PIT	36	QB	16	16	452	675	67.0	5129	34	16	7.6	11.3	320.6	96.5	71.0	24
1	2	Andrew Luck*	IND	29	QB	16	16	430	639	67.3	4593	39	15	7.2	10.7	287.1	98.7	69.4	18
2	3	Matt Ryan	ATL	33	QB	16	16	422	608	69.4	4924	35	7	8.1	11.7	307.8	108.1	68.5	42
3	4	Kirk Cousins	MIN	30	QB	16	16	425	606	70.1	4298	30	10	7.1	10.1	268.6	99.7	58.2	40
4	5	Aaron Rodgers*	GNB	35	QB	16	16	372	597	62.3	4442	25	2	7.4	11.9	277.6	97.6	54.4	49

```
removed_NaN.shape
```

```
(69, 19)
```

# Make sure to restart the index!

---

```
removed_NaN.loc[50:60]
```

	Rk	Player	Tm	Age	Pos	G	GS	Cmp	Att	Cmp%	Yds	TD	Int	Y/A	Y/C	Y/G	Rate	QBR	Sk
51	52	Mark Sanchez	WAS	32	qb	2	1	19	35	54.3	138	0	3	3.9	7.3	69.0	28.0	4.4	7
52	53	Kyle Allen	CAR	22	qb	2	1	20	31	64.5	266	2	0	8.6	13.3	133.0	113.1	96.4	0
53	54	Matt Barkley	BUF	28	qb	1	1	15	25	60.0	232	2	0	9.3	15.5	232.0	117.4	83.4	1
54	55	Teddy Bridgewater	NOR	26	qb	5	1	14	23	60.9	118	1	1	5.1	8.4	23.6	70.6	39.8	2
58	59	Taysom Hill	NOR	28	te/wr	16	4	3	7	42.9	64	0	1	9.1	21.3	4.0	36.3	41.1	1

# Make sure to restart the index!

```
removed_NaN = removed_NaN.reset_index(drop=True)  
removed_NaN.loc[50:60]
```

	Rk	Player	Tm	Age	Pos	G	GS	Cmp	Att	Cmp%	Yds	TD	Int	Y/A	Y/C	Y/G	Rate	QBR	Sk
50	52	Mark Sanchez	WAS	32	qb	2	1	19	35	54.3	138	0	3	3.9	7.3	69.0	28.0	4.4	7
51	53	Kyle Allen	CAR	22	qb	2	1	20	31	64.5	266	2	0	8.6	13.3	133.0	113.1	96.4	0
52	54	Matt Barkley	BUF	28	qb	1	1	15	25	60.0	232	2	0	9.3	15.5	232.0	117.4	83.4	1
53	55	Teddy Bridgewater	NOR	26	qb	5	1	14	23	60.9	118	1	1	5.1	8.4	23.6	70.6	39.8	2
54	59	Taysom Hill	NOR	28	te/wr	16	4	3	7	42.9	64	0	1	9.1	21.3	4.0	36.3	41.1	1
55	68	Derrick Henry	TEN	24	RB	16	12	2	3	66.7	14	0	0	4.7	7.0	0.9	77.1	59.3	0
56	72	Julian Edelman	NWE	32	WR	12	12	2	2	100.0	43	0	0	21.5	21.5	3.6	118.7	94.3	0
57	74	Jarvis Landry*	CLE	26	WR	16	14	1	2	50.0	63	0	0	31.5	63.0	3.9	95.8	100.0	0
58	75	Mohamed Sanu	ATL	29	WR	16	16	1	2	50.0	5	1	0	2.5	5.0	0.3	95.8	32.3	0
59	77	Nelson Agholor	PHI	25	WR	16	16	1	1	100.0	15	0	0	15.0	15.0	0.9	118.7	100.0	0
60	78	Danny Amendola	MIA	33	WR	15	15	1	1	100.0	28	1	0	28.0	28.0	1.9	158.3	100.0	0



# Option #2, Fill each column differently

---

- The first step in this is to determine which columns have NaN values. This is done pretty easily using the following function

```
# Find out which columns have NaN values  
passing_stats.columns[passing_stats.isna().any()].tolist()
```

```
['Pos', 'Y/C', 'QBR']
```



# Option #2, Fill each column differently

```
# Set up a dictionary to hold the default values to fill NaN values
defaults = {
    "Pos" : "NA",
    "Y/C" : 0,
    "QBR" : 0
}

# Fill these values in the table
cleaned = passing_stats.fillna(defaults)
cleaned.loc[[50]]
```

	Rk	Player	Tm	Age	Pos	G	GS	Cmp	Att	Cmp%	Yds	TD	Int	Y/A	Y/C	Y/G	Rate	QBR	Sk
50	51	DeShone Kizer	GNB	22	NA	3	0	20	42	47.6	187	0	2	4.5	9.4	62.3	40.5	25.2	4

# Final tips for data cleaning

---

- Always be sure to clean your data before doing any analysis
- Be mindful of the values you are removing from the table
  - Make sure you are aware of the consequences of removing or changing values
  - Excluding values that you should not be excluding can have adverse effects on your final analysis

# Getting information from the table

---

- The next section of this module will focus on trends and visualizations
  - Not designed to be exhaustive, but rather point you in the right direction
- A lot of analysis is designed for you to determine what is important and what isn't
- Think about your goals before you start cleaning and visualizing!

# Let's transition to a new dataset

- Every game log of the 2018 season for every player
- Open up a new jupyter notebook and read in the new file

```
import pandas as pd
```

```
df = pd.read_excel("nba_stats.xlsx")
```

```
df.head()
```

	Name	Player_Game_Number	Date	Age	Team	Location	Opponent	Result	GS	TP	...	ORB	DRB	TRB	AST	STL	BLK	TOV	PF	PTS	Plus
0	Álex Abrines	1	2018-10-16	25-076	OKC	Away	GSW	L (-8)	0	23:28	...	0	2	2	0	0	0	0	2	8	
1	Álex Abrines	2	2018-10-19	25-079	OKC	Away	LAC	L (-16)	0	32:06	...	1	1	2	1	1	0	2	0	10	
2	Álex Abrines	3	2018-10-21	25-081	OKC	Home	SAC	L (-11)	0	5:20	...	0	1	1	2	0	0	0	0	0	
3	Álex Abrines	4	2018-10-25	25-085	OKC	Home	BOS	L (-6)	0	18:33	...	0	1	1	1	0	0	0	2	6	
4	Álex Abrines	5	2018-10-28	25-088	OKC	Home	PHO	W (+7)	0	24:20	...	0	1	1	0	2	0	0	2	2	

# Let's understand the dataset

---

- What do the columns mean?
- What are the types of each column?
- What is the dataset telling me?
- What possible trends could I visualize?

# Get information about a column

```
df["AST"].max()
```

24

```
df["AST"].describe()
```

```
count    26101.000000
mean         2.317268
std         2.525151
min          0.000000
25%          0.000000
50%          2.000000
75%          3.000000
max         24.000000
Name: AST, dtype: float64
```

```
df[df["AST"] == df["AST"].max()]
```

	Name	Player_Game_Number	Date	Age	Team	Location	Opponent	Result	GS	TP	...	ORB	DRB	TRB	AST	STL	BLK	TOV	PF	PT
24962	Russell Westbrook	33	2019-01-10	30-059	OKC	Away	SAS	L (-7)	1	49:35	...	2	11	13	24	2	0	3	4	2

1 rows × 29 columns

# Plotting


---

- At the top of your jupyter notebook, we will need to import a new library that allows you to plot data
  - `import matplotlib.pyplot as plt`
- This is not included in the pandas library
- Why are data visualizations important?
- Good tutorial:  
<http://queirozf.com/entries/pandas-dataframe-plot-examples-with-matplotlib-pyplot>

# Bar Graphs

- Provide an x axis, y axis, and the type of graph you want to see
  - In this case, the kind is “bar”
- What are bar graphs good for?
  - This example might not be the best idea for a bar graph
  - How can we make it better?

```
# This plot shows how many points he scored in each game  
# It looks super congested!  
plt = harden.plot(x="Date", y="PTS", kind="bar")
```





# Bar Graph Part 2, Creating a smaller range

---

- Looking at every date in the season was too much. How about we narrow down our range to be smaller?
- Let's look at the first two weeks of December (arbitrary)
- Is there a way we can easily take our DataFrame to slice out this region?
  - Because the dates are stored as strings, it makes it difficult to compare using inequalities
  - How do we know if a date is in that range?
- We know the dates that we want are in the format: YYYY-MM-DD
  - We also know we want the dates of 2018-12-01 to 2018-12-14
  - How can we do this? For loop!

# What is missing here?

---

```
# Let's look at a specific range of dates (first two weeks of decem
# Date format in the table is YYYY-MM-DD
dates = []
for day in range(14):
    date = "2018-12-" + str(day)
    dates.append(date)
print(dates)
```

```
['2018-12-0', '2018-12-1', '2018-12-2', '2018-12-3', '2018-12-4',
'2018-12-5', '2018-12-6', '2018-12-7', '2018-12-8', '2018-12-9',
'2018-12-10', '2018-12-11', '2018-12-12', '2018-12-13']
```

# Fixed Code

---

```
# Lets fix the code in the previous cell
dates = []
for day in range(14):
    if day < 10:
        str_day = "0" + str(day + 1)
    else:
        str_day = str(day + 1)
    date = "2018-12-" + str_day
    dates.append(date)
print(dates)
```


```
['2018-12-01', '2018-12-02', '2018-12-03', '2018-12-04', '2018-12-05', '2018-12-06', '2018-12-07', '2018-12-08', '2018-12-09', '2018-12-10', '2018-12-11', '2018-12-12', '2018-12-13', '2018-12-14']
```

# Fixed Bar Graph

```
# Lets replot our histogram
small_harden = harden[harden["Date"].isin(dates)]
small_harden
```

	Name	Player_Game_Number	Date	Age	Team	Location	Opponent	Result	GS	TP
18	James Harden	19	2018-12-01	29-097	HOU	Home	CHI	W (+16)	1	30:22
19	James Harden	20	2018-12-03	29-099	HOU	Away	MIN	L (-12)	1	37:01
20	James Harden	21	2018-12-06	29-102	HOU	Away	UTA	L (-27)	1	27:58
21	James Harden	22	2018-12-08	29-104	HOU	Away	DAL	L (-3)	1	36:39
22	James Harden	23	2018-12-11	29-107	HOU	Home	POR	W (+7)	1	32:28
23	James Harden	24	2018-12-13	29-109	HOU	Home	LAL	W (+15)	1	35:28

```
plt = small_harden.plot(x="Date", y="PTS", kind="bar")
```




# Histograms

---

- Histograms look similar to bar graphs but show groupings of data instead of individual values
  - When creating a histogram, determine the column you want to analyze and the “bins” you want to use
- What are histograms good for?
  - Determining which values are the most common
  - Examining distributions


```
plt = allstars_df[["PTS"]].plot(  
    kind="hist",  
    bins = [0, 5, 10, 15, 20, 25, 30, 35, 40, 45, 50, 55, 60, 65],  
    rwidth=0.8  
)
```



# Same Histogram, different bin size!

---

```
# Histogram displaying frequency of scoring grouped in bins of 10  
plt = allstars_df[["PTS"]].plot(  
    kind="hist",  
    bins = [0, 10, 20, 30, 40, 50, 60, 70],  
    rwidth=0.8  
)
```




# Scatter Plots

---

- Scatter plots attempt to display the relationship between two variables
  - Why is this useful?
- Provide two variables you want to see and plot!

```
# Scatter plot between Points and Assists  
plt = allstars_df.plot(x="PTS", y = "AST", kind="scatter")
```



# Further applications for you to explore

---

- Improving your visualizations
  - Adding titles, axis labels
  - Using different colors
  - Exploring different types of visualizations
- Linear regression
  - Too long for this tutorial
- Graphing multiple trends on the same plot
  - Use different colors to display multiple trends



# Any Questions?

---

# What do you want to see in the last module?

---

## Options

- Webscraping live demo
- Linear regression tutorial
- More visualization help
- General live coding demo
- Any other options?